

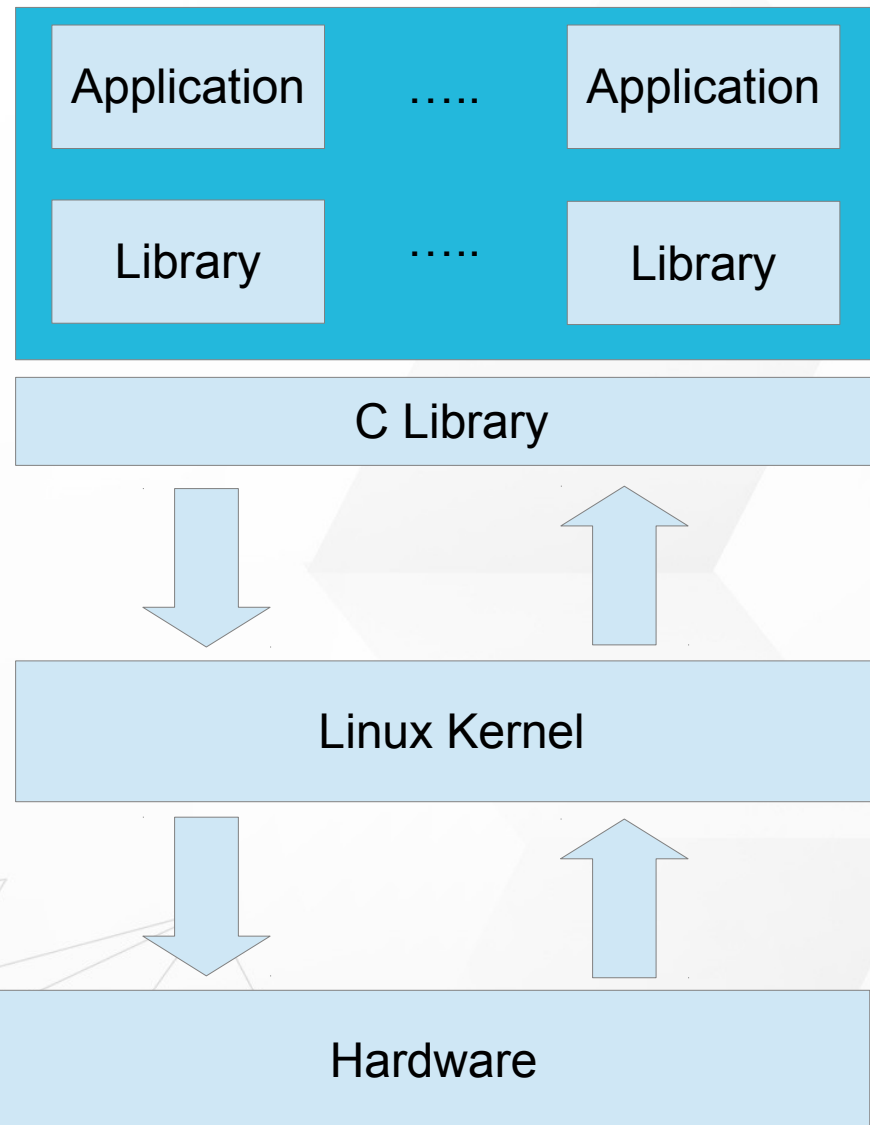
Linux kernel



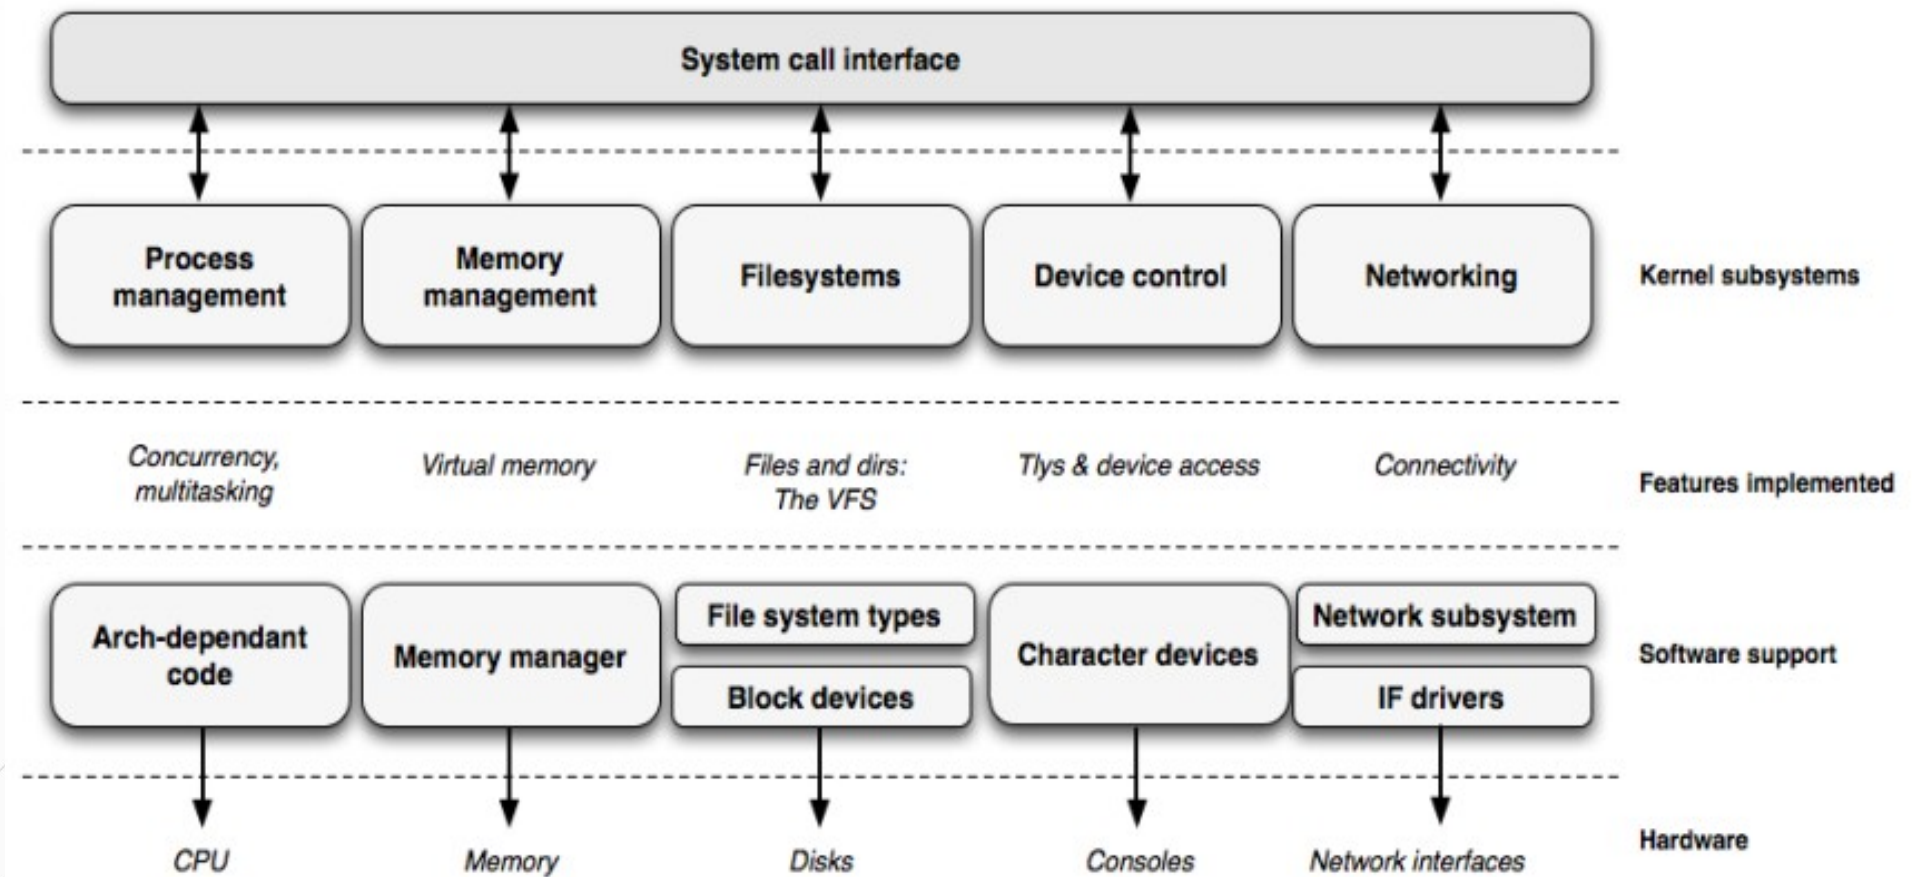
Linux kernel key features

- Portability and hardware support
- Scalability
- Compliance to standards and interoperability
- Exhaustive networking support
- Stability and reliability
- Modularity
- Easy to program.

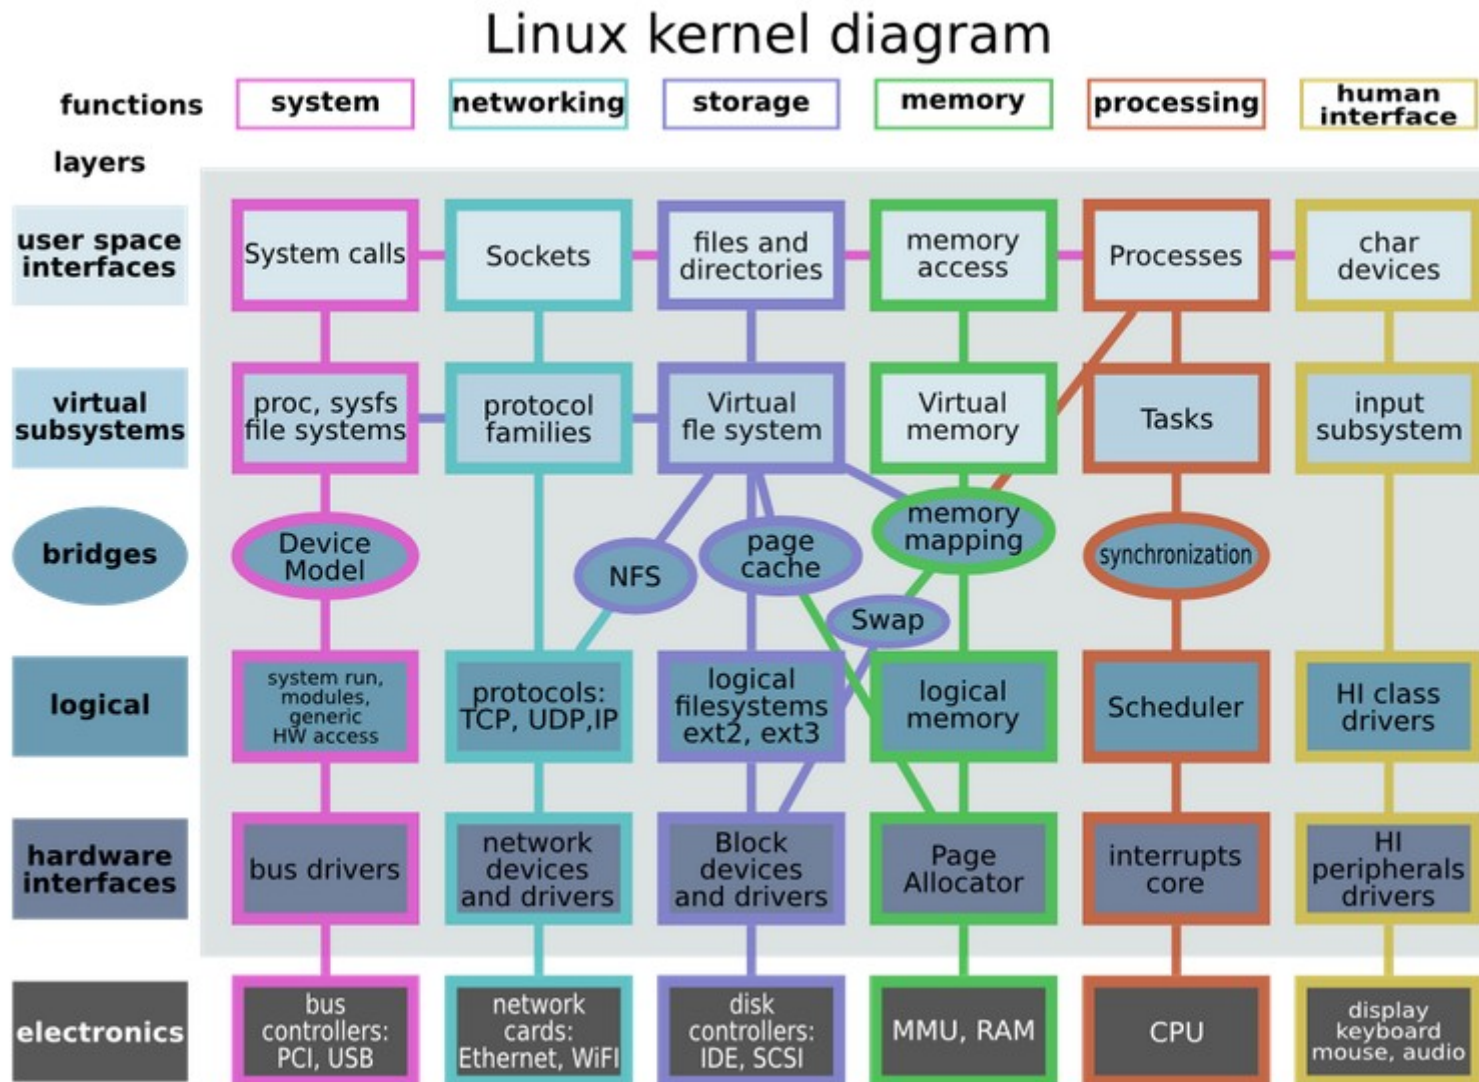
Linux kernel in the system



Linux kernel



Linux kernel





Kernel Source

➤ <https://www.kernel.org/>

➤ Many chip vendors

➤ kernel sub-communities

➤ Architecture communities

- ARM, MIPS, PowerPC ...

➤ device drivers communities

- I2C, SPI, USB, PCI, network ...



Programming language

- Implemented in C like all Unix systems
- A little Assembly is used too
- **No C++ used**
- **No C library**
- **No floating point computation**
- Kernel code has to supply its own library implementations
 - X : printf(), memset(), malloc(),...
 - O: printk(), memset(), kmalloc()



Linux sources structure

- arch/<ARCH>

- Architecture specific code

 - arch/<ARCH>/mach-<machine>

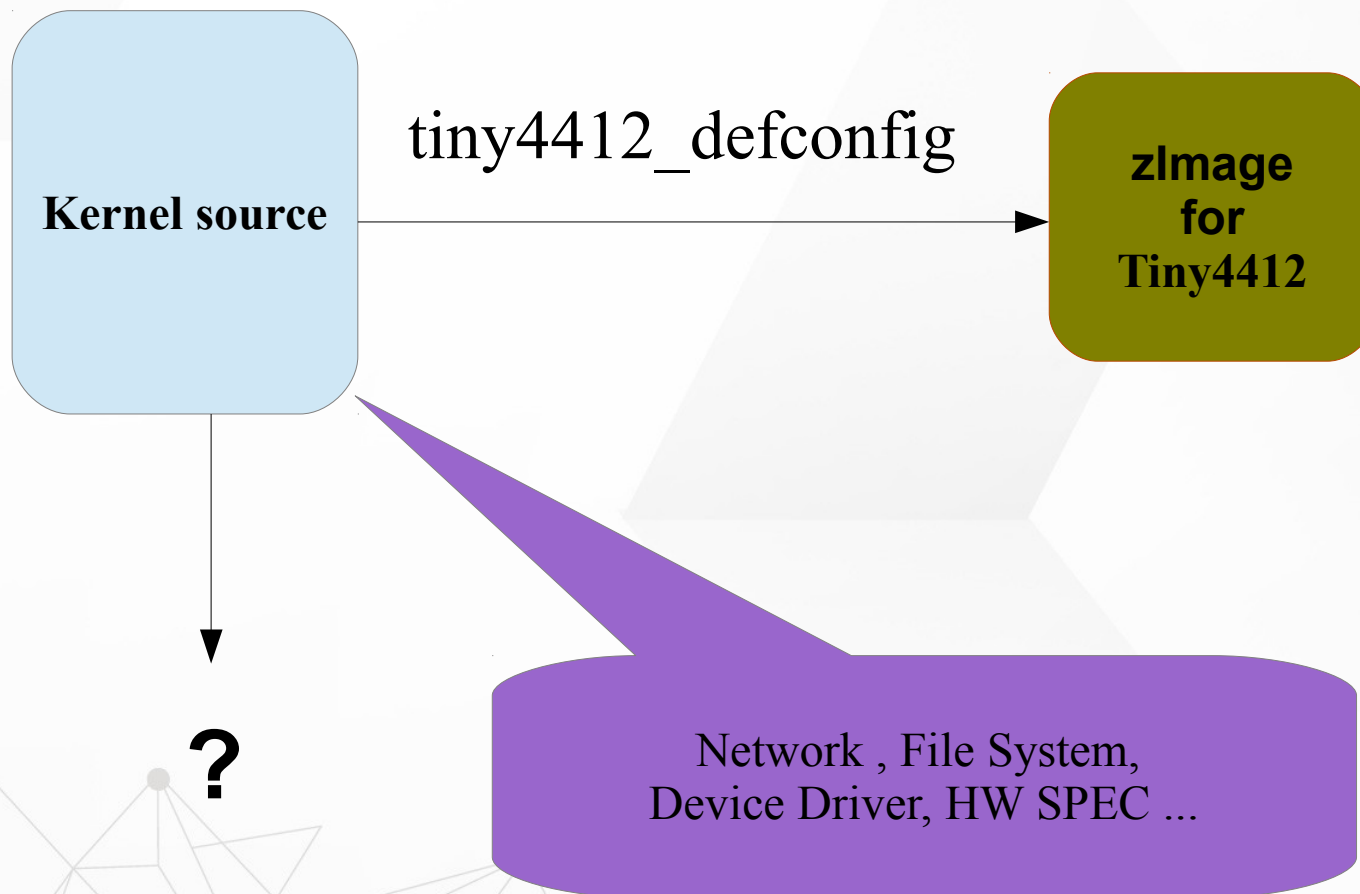
 - arch/<ARCH>/include/asm

 - arch/<ARCH>/boot/

- driver/

- Documentation/

Kernel configuration





Kernel configuration

- The kernel configuration and build system is based on multiple Makefiles
- The configuration is stored in the **.config** file at the root of kernel sources
- As options have dependencies, typically never edited by hand, but through graphical or text interfaces
 - **make menuconfig** → Text
 - **make xconfig** → graphical

Kernel configuration

```
slash@slash-huang:linux_3.5.0_tiny4412$ ls -la
total 574796
drwxr-xr-x  24 slash slash      4096 Sep  9 14:14 .
drwxrwxr-x   4 slash slash      4096 Sep  8 11:44 ..
drwxr-xr-x  29 slash slash      4096 Apr 22  2014 arch
drwxr-xr-x   3 slash slash      4096 Sep  9 14:06 block
-rwxrwxr-x   1 slash slash        37 Mar  4  2015 change_lib.sh
-rw-rw-r--   1 slash slash    83180 Sep  9 14:04 .config
-rw-r--r--   1 slash slash   18693 Apr 22  2014 COPYING
-rw-r--r--   1 slash slash   94956 Apr 22  2014 CREDITS
drwxr-xr-x   3 slash slash      4096 Sep  9 14:12 crypto
```

.config

\$KERNEL_PATH/arch/arm/configs/
tiny4412_defconfig

```
slash@slash-huang:linux_3.5.0_tiny4412$ ls arch/arm/configs/
acs5k_defconfig      cm_x2xx_defconfig    g3evm_defconfig      lpd270_defconfig     pcm027_defconfig     spear13xx_defconfig
acs5k_tiny_defconfig cm_x300_defconfig    g4evm_defconfig      lubbock_defconfig    pleb_defconfig       spear3xx_defconfig
afeb9260_defconfig   cns3420vb_defconfig  h3600_defconfig      mackerel_defconfig   pnx4008_defconfig    spear6xx_defconfig
ag5evm_defconfig     colibri_pxa270_defconfig h5000_defconfig      magician_defconfig   prima2_defconfig     spitz_defconfig
am200epdkit_defconfig colibri_pxa300_defconfig h7201_defconfig      mainstone_defconfig  pxa168_defconfig     stamp9g20_defconfig
ap4evb_defconfig     collie_defconfig     h7202_defconfig      marzen_defconfig     pxa255-idp_defconfig tct_hammer_defconfig
armadillo800eva_defconfig corgi_defconfig      hackkit_defconfig    mini2440_defconfig   pxa3xx_defconfig     tegra_defconfig
assabet_defconfig    cpu9260_defconfig    imote2_defconfig     mmp2_defconfig       pxa910_defconfig     tiny4412_android_defconfig
at91_dt_defconfig    cpu9g20_defconfig    imx_v4_v5_defconfig  msm_defconfig        qil-a9260_defconfig  tiny4412_linux_defconfig
at91rm9200_defconfig da8xx_omapl_defconfig imx_v6_v7_defconfig  mv78xx0_defconfig    raumfeld_defconfig   tiny4412_ubuntu_defconfig
at91sam9260_defconfig davinci_all_defconfig integrator_defconfig mxs_defconfig         realview_defconfig   trizeps4_defconfig
at91sam9261_defconfig dove_defconfig       iop13xx_defconfig    neponset_defconfig   realview-smp_defconfig u300_defconfig
at91sam9263_defconfig ebsa110_defconfig    iop32x_defconfig     netwinder_defconfig  rpc_defconfig         u8500_defconfig
at91sam9g20_defconfig edb7211_defconfig    iop33x_defconfig     netx_defconfig       s3c2410_defconfig    usb-a9260_defconfig
at91sam9g45_defconfig em_x270_defconfig    ixp4xx_defconfig     nhk8815_defconfig    s3c6400_defconfig    versatile_defconfig
at91sam9rl_defconfig ep93xx_defconfig     jornada720_defconfig kirkwood_defconfig   nuc910_defconfig     s5p64x0_defconfig    vexpress_defconfig
at91x40_defconfig    eseries_pxa_defconfig kirkwood_defconfig   nuc950_defconfig     nuc950_defconfig     s5pc100_defconfig    viper_defconfig
badge4_defconfig     exynos4_android_defconfig ks8695_defconfig     omap1_defconfig      omap2plus_defconfig  sam9_l9260_defconfig xcep_defconfig
bcmring_defconfig    ezx_defconfig        kzm9g_defconfig      lart_defconfig       orion5x_defconfig    shannon_defconfig    zeus_defconfig
bonito_defconfig     footbridge_defconfig lpc32xx_defconfig    lart_defconfig       orion5x_defconfig    shark_defconfig      simpad_defconfig
cam60_defconfig      fortunet_defconfig   lpc32xx_defconfig    lart_defconfig       orion5x_defconfig    shark_defconfig      simpad_defconfig
cerfcube_defconfig   fortunet_defconfig   lpc32xx_defconfig    lart_defconfig       orion5x_defconfig    shark_defconfig      simpad_defconfig
slash@slash-huang:linux_3.5.0_tiny4412$
```



Kernel or module?

- The kernel image is a single file, resulting from the linking of all object files that correspond to features enabled in the configuration
- Some features (device drivers, filesystems, etc.) can however be compiled as modules

menuconfig

```
Linux/arm 3.5.0 Kernel Configuration

Arrow keys navigate the menu. <Enter> selects submenus --->. Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M>
modularizes features. Press <Esc><Esc> to exit, <?> for Help, </> for Search. Legend: [*] built-in [ ] excluded <M> module < > module
capable

[*] Patch physical to virtual translations at runtime
    General setup --->
[*] Enable loadable module support --->
[*] Enable the block layer --->
    System Type --->
[ ] FIQ Mode Serial Debugger
    Bus support --->
    Kernel Features --->
    Boot options --->
    CPU Power Management --->
    Floating point emulation --->
    Userspace binary formats --->
    Power management options --->
[*] Networking support --->
    Device Drivers --->
    File systems --->
    Kernel hacking --->
    Security options --->
-*- Cryptographic API --->
    Library routines --->
---
    Load an Alternate Configuration File
    Save an Alternate Configuration File

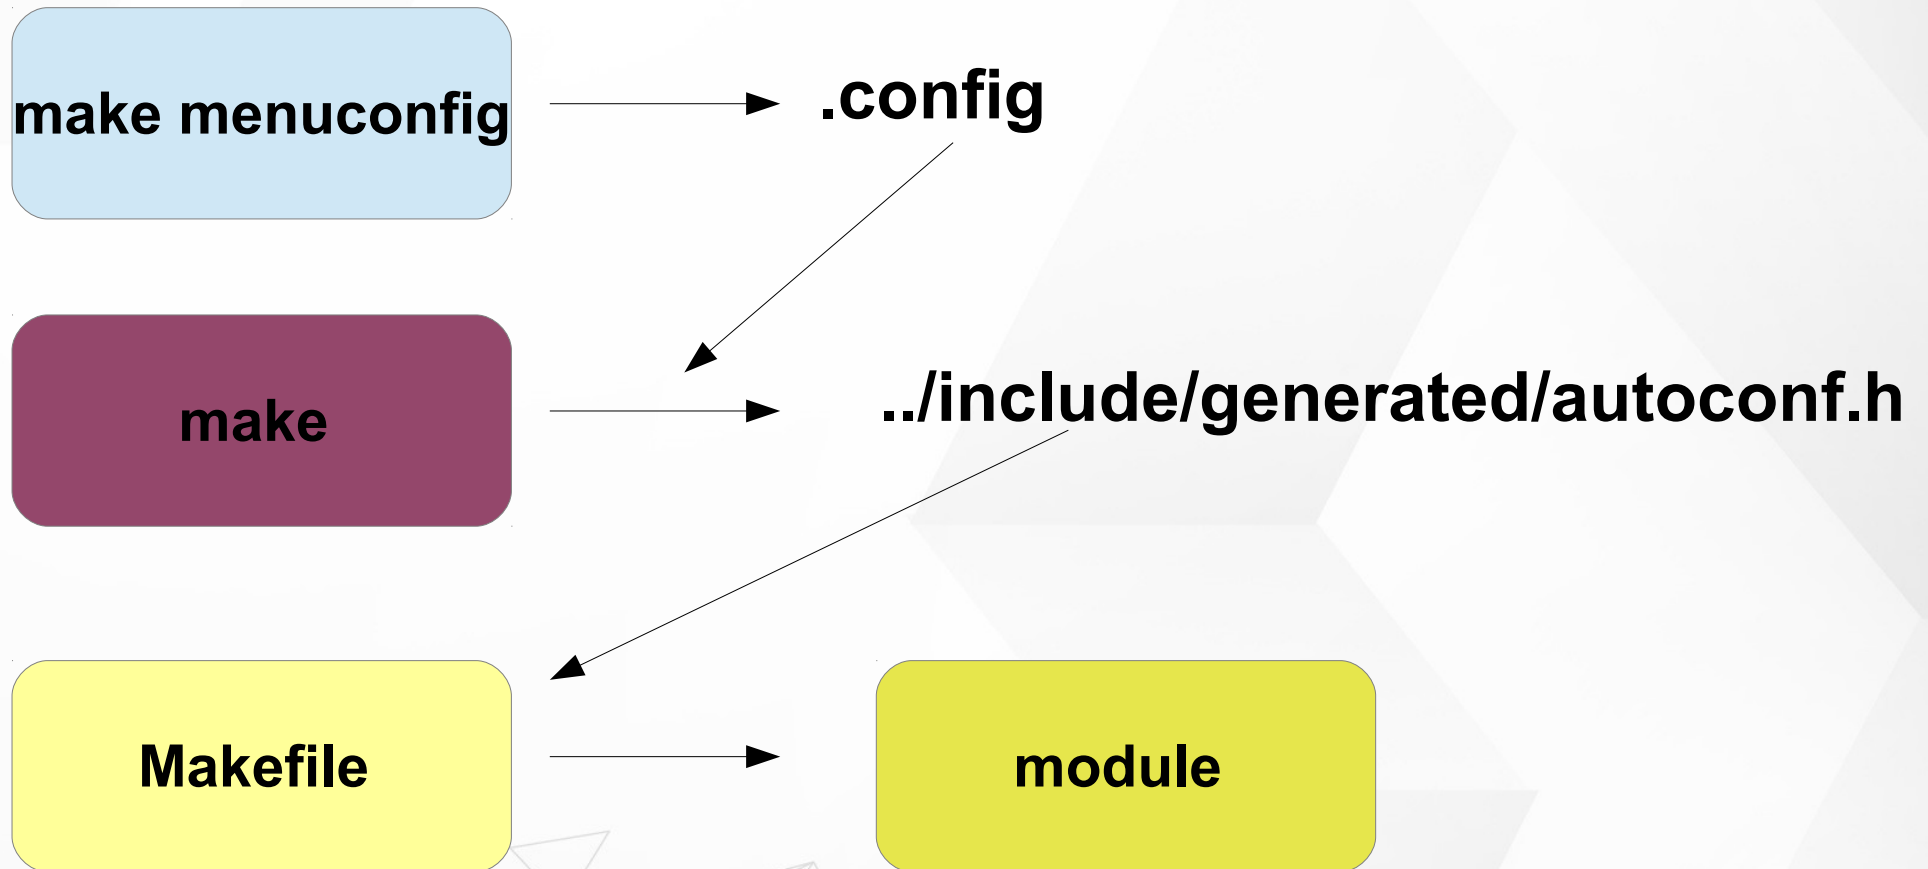
<Select>  < Exit >  < Help >
```


Kernel configuration options

```
^(-)
  PTP clock support --->
  *- GPIO Support --->
  < > Dallas's 1-wire support --->
  [*] Power supply class support --->
  <M> Hardware Monitoring support --->
  <*> Generic Thermal sysfs driver --->
  [*] Watchdog Timer Support --->
      Sonics Silicon Backplane --->
      Broadcom specific AMBA --->
      Multifunction device drivers --->
  [ ] Voltage and Current Regulator Support --->
  <M> Multimedia support --->
      Graphics support --->
  <*> Sound card support --->
      HID support --->
  [*] USB support --->
  <*> MMC/SD/SDIO card support --->
  < > Sony MemoryStick card support (EXPERIMENTAL) --->
  [*] LED Support --->
  <*> Switch class support --->
  [ ] Accessibility support --->
  [*] Real Time Clock --->
  *- DMA Engine support --->
v(+)

<Select>  < Exit >  < Help >
```


Configuration



Corresponding .config file excerpt

```
CONFIG_BLK_DEV_SD=y
CONFIG_CHR_DEV_ST is not set
CONFIG_CHR_DEV_OSST is not set
CONFIG_BLK_DEV_SR is not set
CONFIG_CHR_DEV_SG=y
CONFIG_CHR_DEV_SCH is not set
CONFIG_SCSI_MULTI_LUN=y
CONFIG_SCSI_CONSTANTS is not set
CONFIG_SCSI_LOGGING is not set
CONFIG_SCSI_SCAN_ASYNC is not set
CONFIG_SCSI_WAIT_SCAN=m
```

```
# --- NOTE ORDERING HERE ---
# For kernel non-modular link, transport attributes need to
# be initialised before drivers
# -----
obj-$(CONFIG_SCSI_SPI_ATTRS) += scsi_transport_spi.o
obj-$(CONFIG_SCSI_FC_ATTRS) += scsi_transport_fc.o
obj-$(CONFIG_SCSI_ISCSI_ATTRS) += scsi_transport_iscsi.o
obj-$(CONFIG_SCSI_SAS_ATTRS) += scsi_transport_sas.o
obj-$(CONFIG_SCSI_SAS_LIBSAS) += libsas/
obj-$(CONFIG_SCSI_SRP_ATTRS) += scsi_transport_srp.o
obj-$(CONFIG_SCSI_DH) += device_handler/
```

SCSI Transports

```
CONFIG_SCSI_SPI_ATTRS is not set
CONFIG_SCSI_FC_ATTRS is not set
CONFIG_SCSI_ISCSI_ATTRS is not set
CONFIG_SCSI_SAS_ATTRS is not set
CONFIG_SCSI_SAS_LIBSAS is not set
CONFIG_SCSI_SRP_ATTRS is not set
CONFIG_SCSI_LOWLEVEL=y
CONFIG_ISCSI_TCP is not set
CONFIG_ISCSI_BOOT_SYSFS is not set
```

\$(KERNEL_PATH)/drivers/scsi/Makefile

.config



Kernel cleanup targets

➤ **make clean**

➤ Kernel cleanup targets

➤ **make mrproper**

➤ Remove all generated files (.config)

➤ **make distclean**

➤ Remove editor backup and patch reject files



Specifying cross-compilation

- make **ARCH=arm CROSS_COMPILE=arm-linux- ...**
- export ARCH=arm
export CROSS_COMPILE=arm-linux
- Add above setting to script
 - source set_arm_toolchain.sh

Predefined configuration files

- Default configuration files available, per board or per-CPU family
 - They are stored in **arch/<arch>/configs/**, and are just minimal .config files
- **make tiny4412_linux_defconfig**
- To create your own default configuration file
 - **make savedefconfig**, to create a minimal configuration file
 - **mv defconfig arch/<arch>/configs/myown_defconfig**

Kernel compilation

➤ make

➤ To run multiple jobs in parallel if you have multiple CPU cores

- **make -j4**

➤ Generates

➤ arch/<arch>/boot/*Image

- zImage for ARM,

Make help

```
acs5k_defconfig      - Build for acs5k
acs5k_tiny_defconfig - Build for acs5k_tiny
afeb9260_defconfig   - Build for afeb9260
ag5evm_defconfig     - Build for ag5evm
am200epdkit_defconfig - Build for am200epdkit
ap4evb_defconfig     - Build for ap4evb
armadillo800eva_defconfig - Build for armadillo800eva
assabet_defconfig    - Build for assabet
at91_dt_defconfig    - Build for at91_dt
at91rm9200_defconfig - Build for at91rm9200
at91sam9260_defconfig - Build for at91sam9260
at91sam9261_defconfig - Other generic targets:
at91sam9263_defconfig -   all - Build all targets marked with [*]
at91sam9g20_defconfig - * vmlinux - Build the bare kernel
at91sam9g45_defconfig - * modules - Build all modules
at91sam9rl_defconfig  - modules_install - Install all modules to INSTALL_MOD_PATH (default: /)
at91x40_defconfig     - firmware_install - Install all firmware to INSTALL_FW_PATH
                        - (default: $(INSTALL_MOD_PATH)/lib/firmware)
badge4_defconfig      - dir/ - Build all files in dir and below
bcmring_defconfig     - dir/file.[oisS] - Build specified target only
bonito_defconfig      - dir/file.lst - Build specified mixed source/assembly target only
cam60_defconfig       - (requires a recent binutils and recent build (System.map))
cerfcube_defconfig    - dir/file.ko - Build module including final link
cm_x2xx_defconfig     - modules_prepare - Set up for building external modules
cm_x300_defconfig     - tags/TAGS - Generate tags file for editors
cns3420vb_defconfig   - cscope - Generate cscope index
colibri_pxa270_defconfig - gtags - Generate GNU GLOBAL index
colibri_pxa300_defconfig - kernelrelease - Output the release version string
collie_defconfig      - kernelversion - Output the version stored in Makefile
corgi_defconfig       - headers_install - Install sanitised kernel headers to INSTALL_HDR_PATH
                        - (default: /home/xlloss/work/tiny-4412/build/linux_3.5.0_tiny4412/usr)
cpu9260_defconfig     - Build for cpu9260
cpu9g20_defconfig     - Build for cpu9g20
```

How to zImage

objcopy




Kernel proper




Stripped kernel

gzip






Compressed kernel

as



LD



Piggy.gz + bootstrap

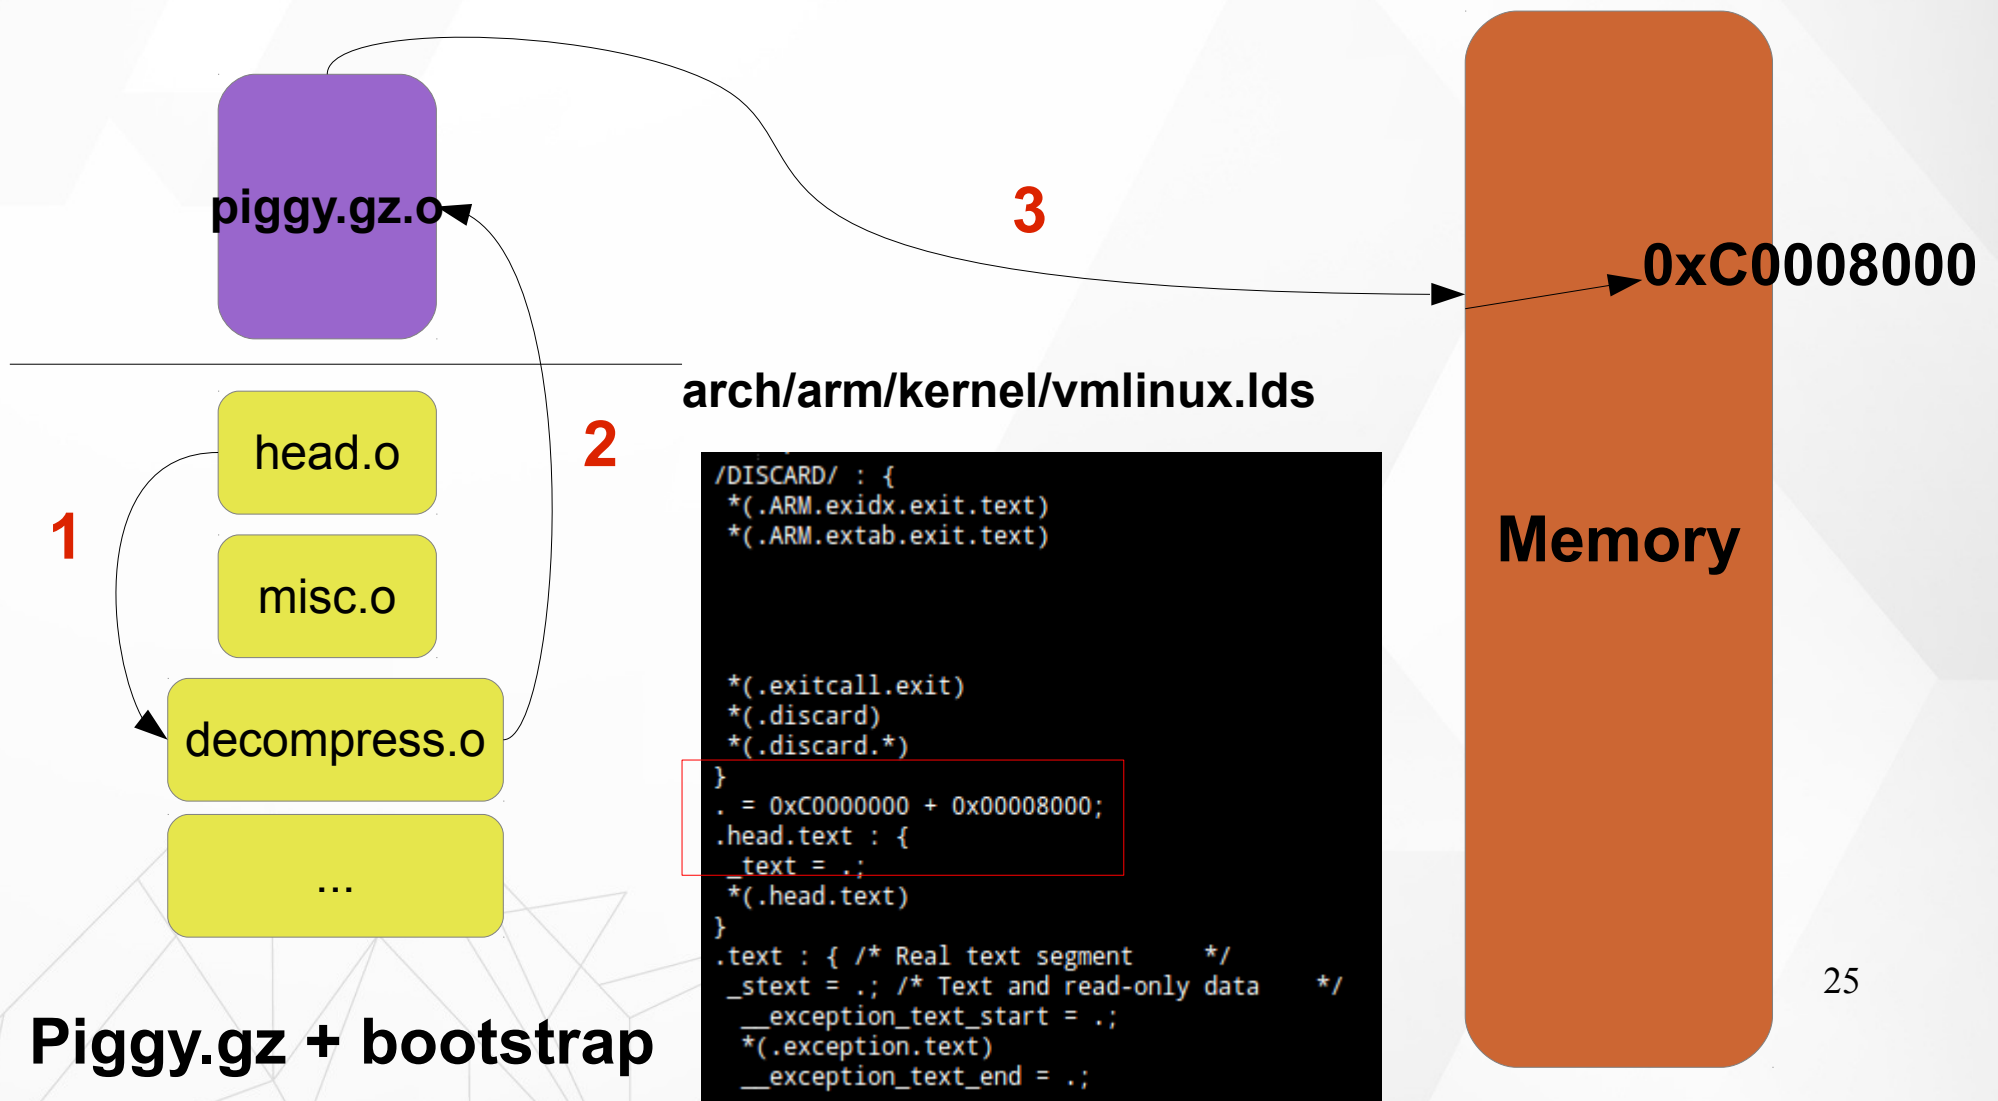


Composite kernel image

objcopy



Kernel boot



Linux boot

`$(linux)/arch/arm/boot/compressed/head.S`

entering

```
/*
 * sort out different calling conventions
 */
.align
.arm                @ Always enter in ARM state
start:|
.type    start,#function
.rept    7
    mov r0, r0
.endr

ARM(     mov r0, r0      )
ARM(     b   1f         )
THUMB(   adr r12, BSYM(1f) )
THUMB(   bx  r12        )

    .word  0x016f2818    @ Magic numbers to help the loader
    .word  start        @ absolute load/run zImage address
    .word  _edata       @ zImage end address
THUMB(   .thumb         )
1:      mov r7, r1       @ save architecture ID
        mov r8, r2       @ save atags pointer

#ifndef __ARM_ARCH_2__
/*
 * Booting from Angel - need to enter SVC mode and disable
 * FIQs/IRQs (numeric definitions from angel arm.h source).
 * We only do this if we were in user mode on entry.
 */
```

Check overwrite

```
/*
 * Check to see if we will overwrite ourselves.
 * r4 = final kernel address
 * r9 = size of decompressed image
 * r10 = end of this image, including bss/stack/malloc space if non XIP
 * We basically want:
 * r4 - 16k page directory >= r10 -> OK
 * r4 + image length <= address of wont_overwrite -> OK
 */
    add r10, r10, #16384
    cmp r4, r10
    bhs wont_overwrite
    add r10, r4, r9
    adr r9, wont_overwrite
    cmp r10, r9
    bls wont_overwrite
*/
```

Relocate

```
/*
 * Relocate ourselves past the end of the decompressed kernel.
 *   r6 = _edata
 *   r10 = end of the decompressed kernel
 * Because we always copy ahead, we need to do it from the end and go
 * backward in case the source and destination overlap.
 */
    /*
     * Bump to the next 256-byte boundary with the size of
     * the relocation code added. This avoids overwriting
     * ourself when the offset is small.
     */
    add r10, r10, #((reloc_code_end - restart + 256) & ~255)
    bic r10, r10, #255

    /* Get start of code we want to copy and align it down. */
    adr r5, restart
    bic r5, r5, #31

    sub r9, r6, r5      @ size to copy
    add r9, r9, #31     @ rounded up to a multiple
    bic r9, r9, #31     @ ... of 32 bytes
    add r6, r9, r5
    add r9, r9, r10

1:    ldmdb    r6!, {r0 - r3, r10 - r12, lr}
    cmp r6, r5
```


Decompress kernel

```
/*
 * The C runtime environment should now be setup sufficiently.
 * Set up some pointers, and start decompressing.
 *  r4 = kernel execution address
 *  r7 = architecture ID
 *  r8 = atags pointer
 */
mov r0, r4
mov r1, sp          @ malloc space above stack
add r2, sp, #0x10000 @ 64k max
mov r3, r7
bl decompress_kernel
bl cache_clean_flush
bl cache_off
mov r0, #0          @ must be zero
mov r1, r7          @ restore architecture number
mov r2, r8          @ restore atags pointer
ARM(    mov pc, r4 ) @ call kernel
THUMB(  bx  r4 )    @ entry point is always ARM

.align 2
.type LC0, #object
LC0:   .word LC0      @ r1
       .word __bss_start @ r2
       .word _end      @ r3
       .word _edata    @ r6
       .word input_data_end - 4 @ r10 (inflated size location)
       .word got_start @ r11
```



Kernel startup entry point

- MMU = off, D-cache = off, I-cache = dont care
- R0 = 0, R1 = machine ID, R2 = atags or dtb pointer
- Machine ID
 - linux/arch/arm/tools/mach-types
- kernel at 0xc0008000

Kernel Startup Entry Point

```
/*
 * Kernel startup entry point.
 * -----
 *
 * This is normally called from the decompressor code. The requirements
 * are: MMU = off, D-cache = off, I-cache = dont care, r0 = 0,
 * r1 = machine nr, r2 = atags or dtb pointer.
 *
 * This code is mostly position independent, so if you link the kernel at
 * 0xc0008000, you call this at __pa(0xc0008000).
 *
 * See linux/arch/arm/tools/mach-types for the complete list of machine
 * numbers for r1.
 *
 * We're trying to keep crap to a minimum; DO NOT add any machine specific
 * crap here - that's what the boot loader (or in extreme, well justified
 * circumstances, zImage) is for.
 */
.arm

__HEAD
ENTRY(stext)

THUMB( adr r9, BSYM(1f) ) @ Kernel is always entered in ARM.
THUMB( bx r9 ) @ If this is a Thumb-2 kernel,
THUMB( .thumb ) @ switch to Thumb now.
THUMB(1:

    setmode PSR_F_BIT | PSR_I_BIT | SVC_MODE, r9 @ ensure svc mode
                                           @ and irq's disabled
    mrc p15, 0, r9, c0, c0 @ get processor id
    bl __lookup_processor_type @ r5=procinfo r9=cpuid
    movs r10, r5 @ invalid processor (r5=0)?
    THUMB( it eq ) @ force fixup-able long branch encoding
```

Machine Init

\$linux/arch/arm/mach-exynos/mach-tiny4412.c

```
        else
            initialize_non_prime_clocks();
    }
#ifdef CONFIG_BUSFREQ_OPP
    dev_add(&busfreq, &exynos4_busfreq.dev);
    ppmu_init(&exynos_ppmu[PPMU_DMC0], &exynos4_busfreq.dev);
    ppmu_init(&exynos_ppmu[PPMU_DMC1], &exynos4_busfreq.dev);
    ppmu_init(&exynos_ppmu[PPMU_CPU], &exynos4_busfreq.dev);
#endif
    set_tmu_platdata();
}

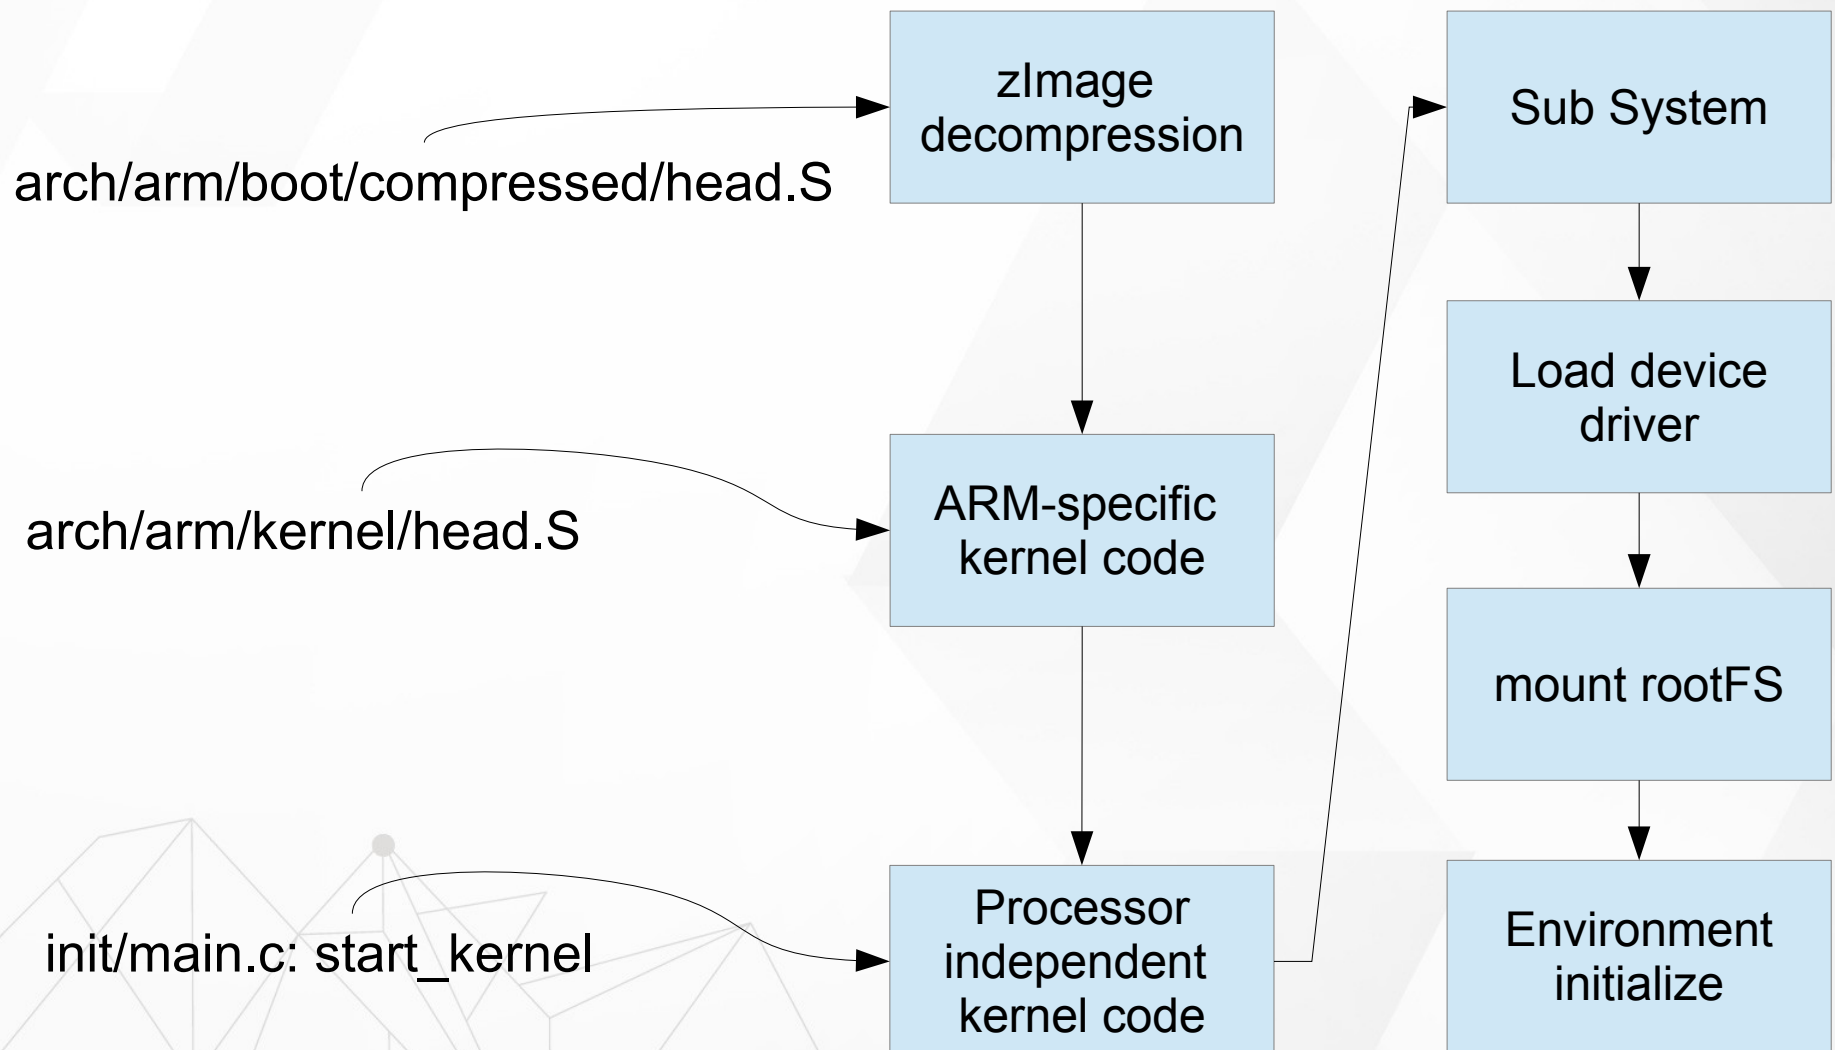
MACHINE_START(TINY4412, "TINY4412")
/* Maintainer: FriendlyARM (www.arm9.net) */
/* Maintainer: Kukjin Kim <kgene.kim@samsung.com> */
/* Maintainer: Changhwan Youn <chaos.youn@samsung.com> */
    .atag_offset    = 0x100,
    .init_irq       = exynos4_init_irq,
    .map_io         = smdk4x12_map_io,
    .handle_irq     = gic_handle_irq,
    .init_machine   = smdk4x12_machine_init,
    .init_late      = exynos_init_late,
    .timer          = &exynos4_timer,
    .restart        = exynos4_restart,
    .reserve        = &smdk4x12_reserve,
MACHINE_END
```

Mach Type

\$linux/arch/arm/tools/mach-types

vybrid_vf4xx	MACH_VYBRID_VF4XX	VYBRID_VF4XX	4148
aria_g25	MACH_ARIA_G25	ARIA_G25	4149
bcm21553	MACH_BCM21553	BCM21553	4150
smdk5410	MACH_SMDK5410	SMDK5410	4151
lpc18xx	MACH_LPC18XX	LPC18XX	4152
oratisparty	MACH_ORATISPARTY	ORATISPARTY	4153
qseven	MACH_QSEVEN	QSEVEN	4154
gmv_generic	MACH_GMV_GENERIC	GMV_GENERIC	4155
th_link_eth	MACH_TH_LINK_ETH	TH_LINK_ETH	4156
tn_muninn	MACH_TN_MUNINN	TN_MUNINN	4157
rampage	MACH_RAMPAGE	RAMPAGE	4158
visstrim_mv10	MACH_VISSTRIM_MV10	VISSTRIM_MV10	4159
mx28_wilma	MACH_MX28_WILMA	MX28_WILMA	4164
msm8625_ffa	MACH_MSM8625_FFA	MSM8625_FFA	4166
vpu101	MACH_VPU101	VPU101	4167
baileys	MACH_BAILEYS	BAILEYS	4169
familybox	MACH_FAMILYBOX	FAMILYBOX	4170
ensemble_mx35	MACH_ENSEMBLE_MX35	ENSEMBLE_MX35	4171
sc_sps_1	MACH_SC_SPS_1	SC_SPS_1	4172
tiny4412	MACH_TINY4412	TINY4412	4608

Linux Kernel Booting





Write rootFS to eMMC

Step 1 – Creation of the actual EXT4.img

```
# dd if=/dev/zero of=rootfs_ext4.img bs=512k count=60
```

Translation of the terms,
bs =blocksize,
count=60, the number of block`s, in our case will result an image of 30 Mb.

To get the exact size of the image that you create use simple maths.

$$60 * 512K = 31457280 \text{ byte} = 30M \text{ bytes}$$

Step 2 Formating the rootfs_ext4.img with EXT4

```
# mkfs.ext4 rootfs_ext4.img
```

It will be a question where you will select yes (Y)

Step 3 mount the directories that we previous created.

```
# mkdir rootfs_ext4 && mount -o loop rootfs_ext4.img rootfs_ext4/
```



Write rootFS to eMMC

Step 4 copy the content from the old system.img in the system_new.img

```
# cp -v -r -p /home/cadtc/tiny4412/experiment/root_mkfs/* ./rootfs_ext4
```

Step 5 sync the files

```
# sync
```

Step 6 Unmounting the partitons.

```
# umount rootfs_ext4/
```

Step 7 reboot EVB and into Linux console with NFS

Step 8 dd rootfs_ext4.img to /dev/mmcblk0p2

```
# dd if=/tmp/rootfs_ext4.img of=/dev/mmcblk0p2
```

Step 9 Modify u-boot args then reboot

```
# setenv bootargs 'noinitrd init=/linuxrc root=/dev/mmcblk0p2 rw  
noinitrd rootfstype=ext4 console=ttySAC0 lcd=S70'
```